

Learning Where and When to Reason in Neuro-Symbolic Inference

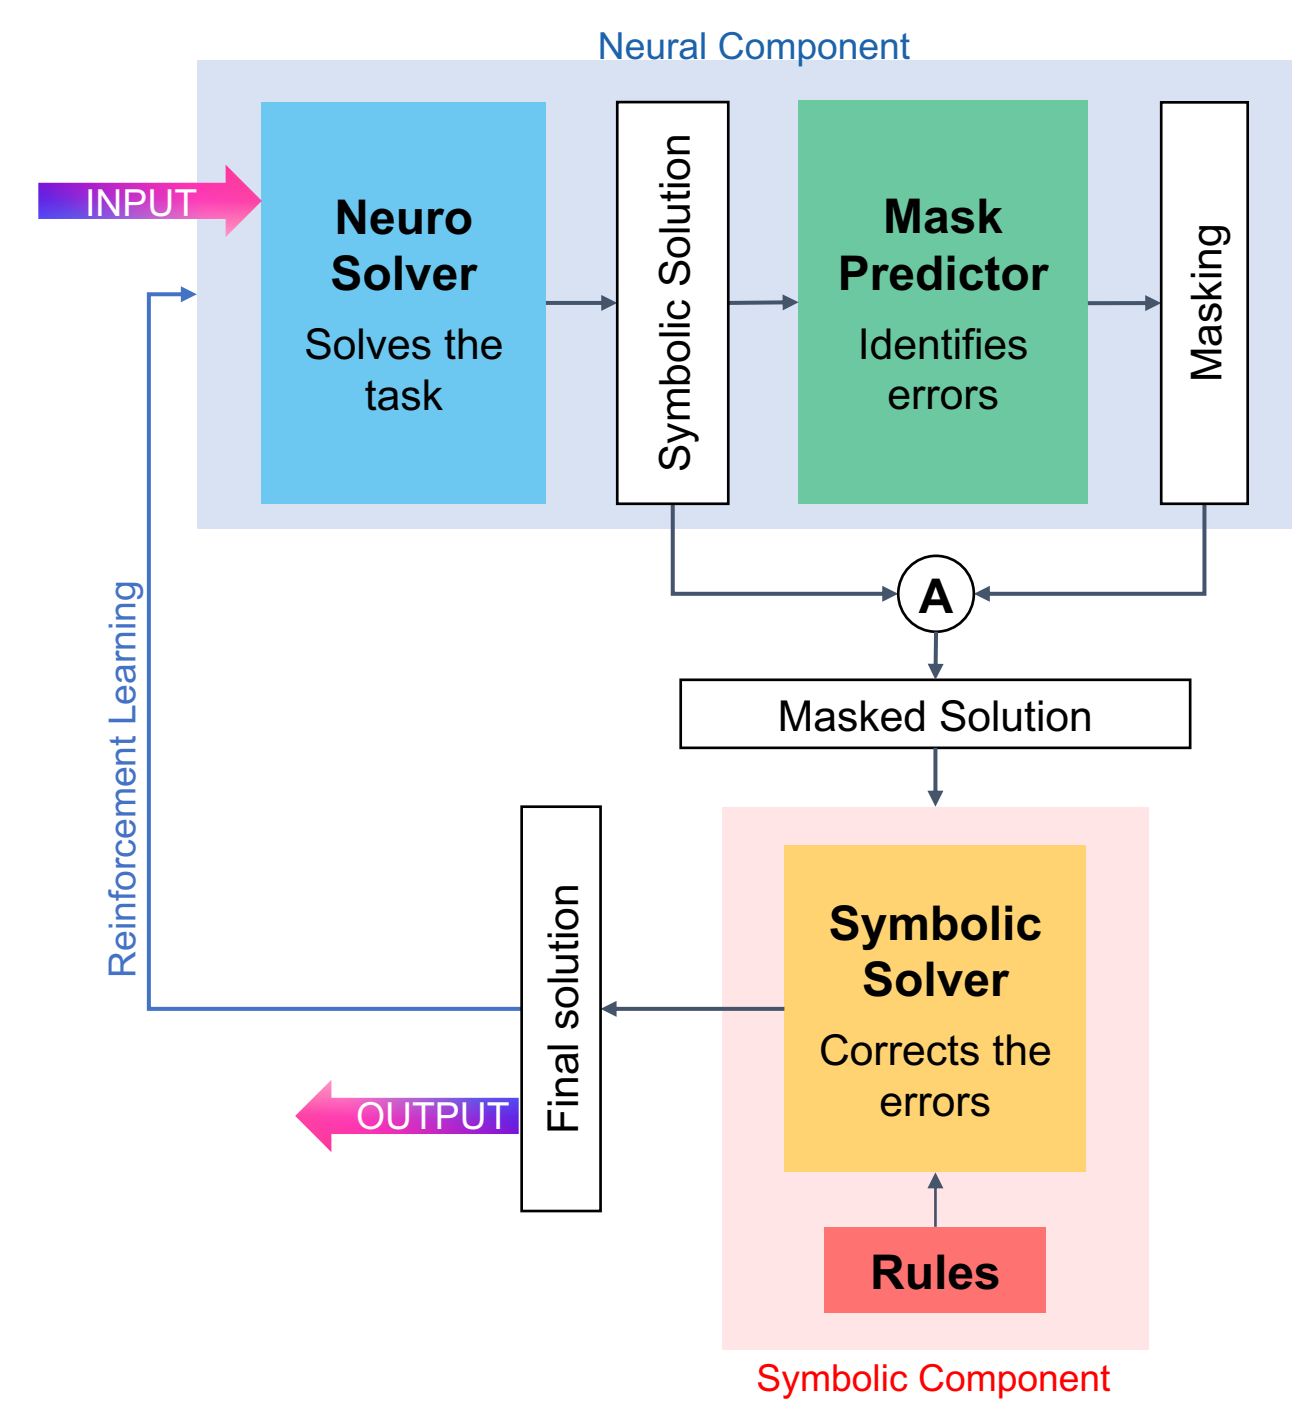
Cristina Cornelio, Jan Stühmer, Shell Xu Hu, Timothy Hospedales ~ Samsung AI Center-Cambridge
ICLR | May 1-5, 2023

Overview

Motivation

- **SOTA**: “Weak”-constraints = enforced only at training time
- **Goal**: Imposition of hard constrains at testing to ensure that the domain-specific knowledge is respected by the predictions
- **Idea**: Neuro-Symbolic pipeline **NASR** (Neural Attention for Symbolic Reasoning)

Architecture



Given: a task to solve & a set of rules \mathcal{R}

1. **Neuro-Solver**: outputs an approximate solution
2. **Mask-Predictor**: identifies the components of the symbolic-solution that do not satisfy the rules \mathcal{R}
3. **Adapter function**: combines the symbolic-solution and the masking to form the masked solution (matching the Symbolic-Solver format)
4. **Symbolic-Solver**: uses the rules \mathcal{R} to correct the masked components of the symbolic solution

- **Symbolic-Solver corrects the Neuro-Solver prediction errors identified by the Mask-Predictor**
- Symbolic reasoning is **not feasible** in many scenarios
- **Mask predictor**: makes the reasoning more efficient, directing the reasoning focus

NASR without RL

Neuro-Solver and **Mask-Predictor** are trained individually (**supervised learning**) and then integrated together

NASR with RL

NASR is then refined using **Reinforcement learning**
 $\mathcal{L}(x; \theta) = -r / \log P_{\theta}(m|ns(x))$

- \mathcal{X} is the set of all possible inputs for the task under consideration
- \mathcal{Y} is the set of all the possible complete solutions
- $\mathcal{Z} = [0, 1]^k$ (where k is the dimension of $y \in \mathcal{Y}$) where 0 indicate a masked element and 1 a non-masked element
- \mathcal{Y}' is equal to \mathcal{Y} with an additional token class 0, corresponding to a masked solution element

Final Hypothesis function maps \mathcal{X} to a probability distribution over \mathcal{Y}

$$f_{\theta}(x) = sb \left(adapt \left(ns(x), argmax \left(mp \left(ns(x) \right) \right) \right) \right) \mathcal{R}$$

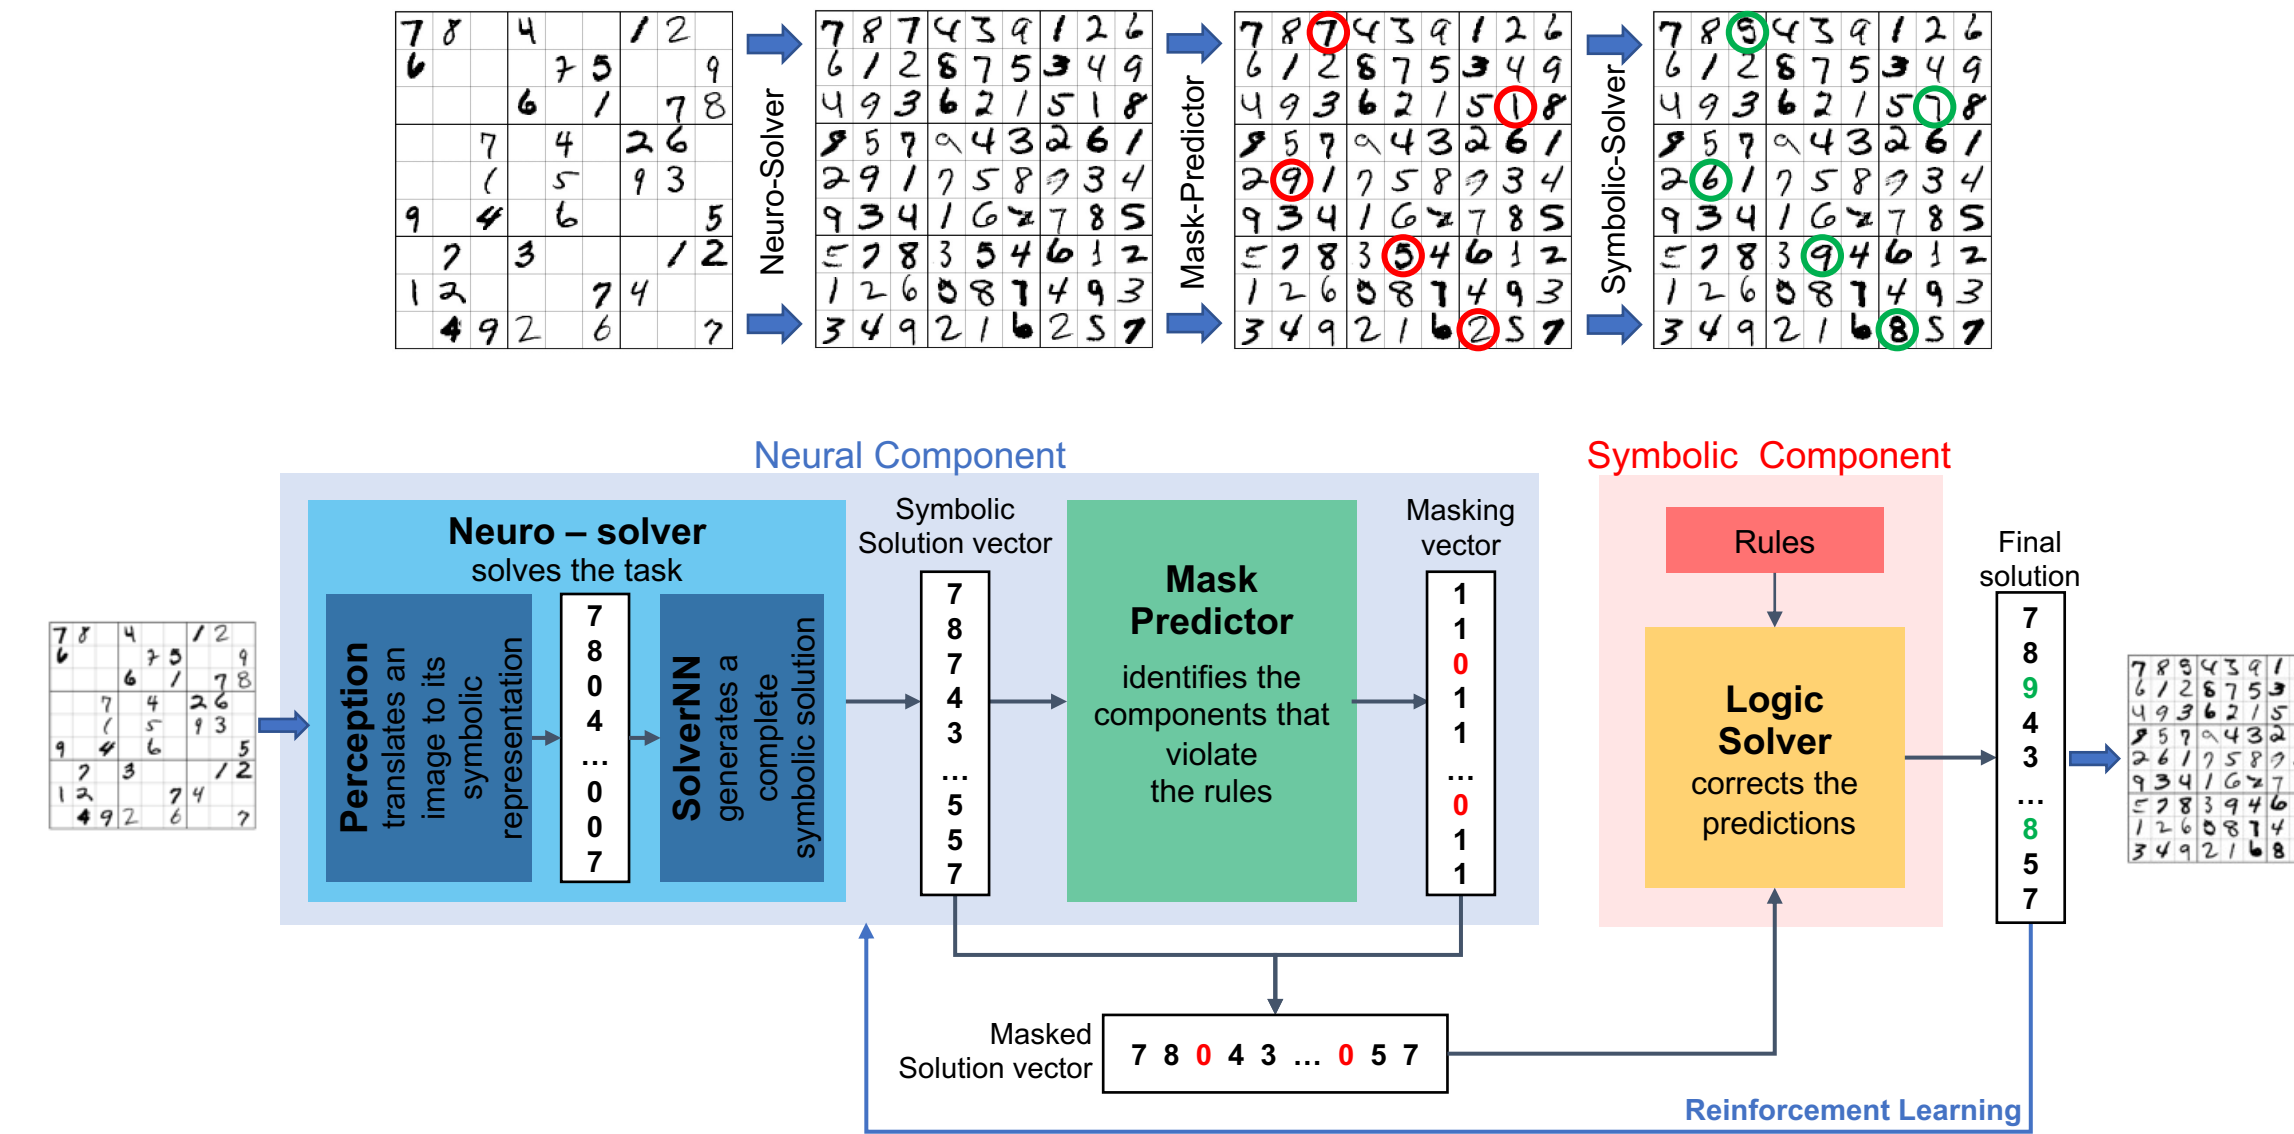
Symbolic Solver maps \mathcal{Y}' to a probability distribution over \mathcal{Y} .

Neuro-Solver maps \mathcal{X} to probability distribution over \mathcal{Y}

Adapter function combines a probability distribution over \mathcal{Y} with an element in \mathcal{Z} (e.g. element wise product)

Mask-Predictor that takes in input a probability distribution over \mathcal{Y} and produce as output a probability distribution over \mathcal{Z}

Results – Visual Sudoku



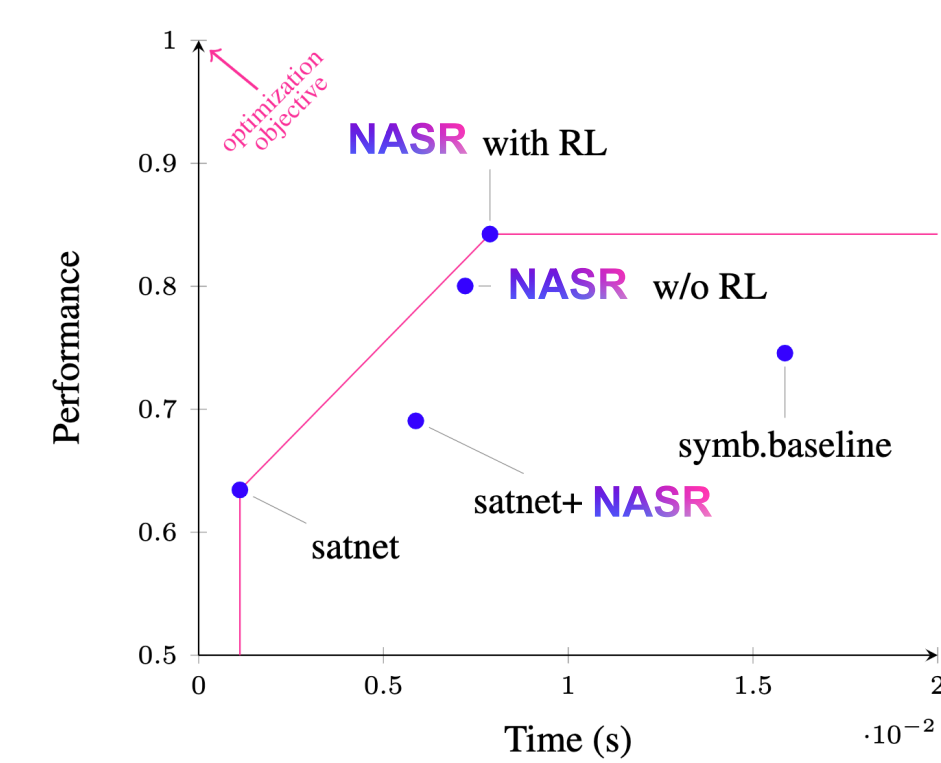
Results summary

- We significantly **outperform the baseline** in most of the cases (and never perform worst);
- We **improve** the performance of an **existing method**, by **integrating it** in our pipeline;
- We are more **robust to noise** compared to the symbolic baseline.


	big kaggle	minimal 17	multiple sol	satnet data
Symbolic baseline	74.56	87.70	63.50	63.20
NeurASP	timeout	89.00*	timeout	timeout
SatNet	63.44	0.00	0.00	60.10
SatNet + NASR (our)	69.05	0.02	24.20	81.40
NASR (our)	84.24	87.00	73.00	82.20

Number of completely correct sudoku boards

Efficiency



Pareto front performance vs. computational time




Performance limiting the computational time

Efficiency = trade-off between:

- performance (percentage of completely correct boards)
- computational time

Attention Maps

When considering the average attention in the transformer of a cell: Focus on the **row**, the **column** and the **3×3 block** (the 3 Sudoku rules)




Cell(1,1)

Cell(5,4)

Cell(8,8)

Mask-Predictor



Cell(1,1)

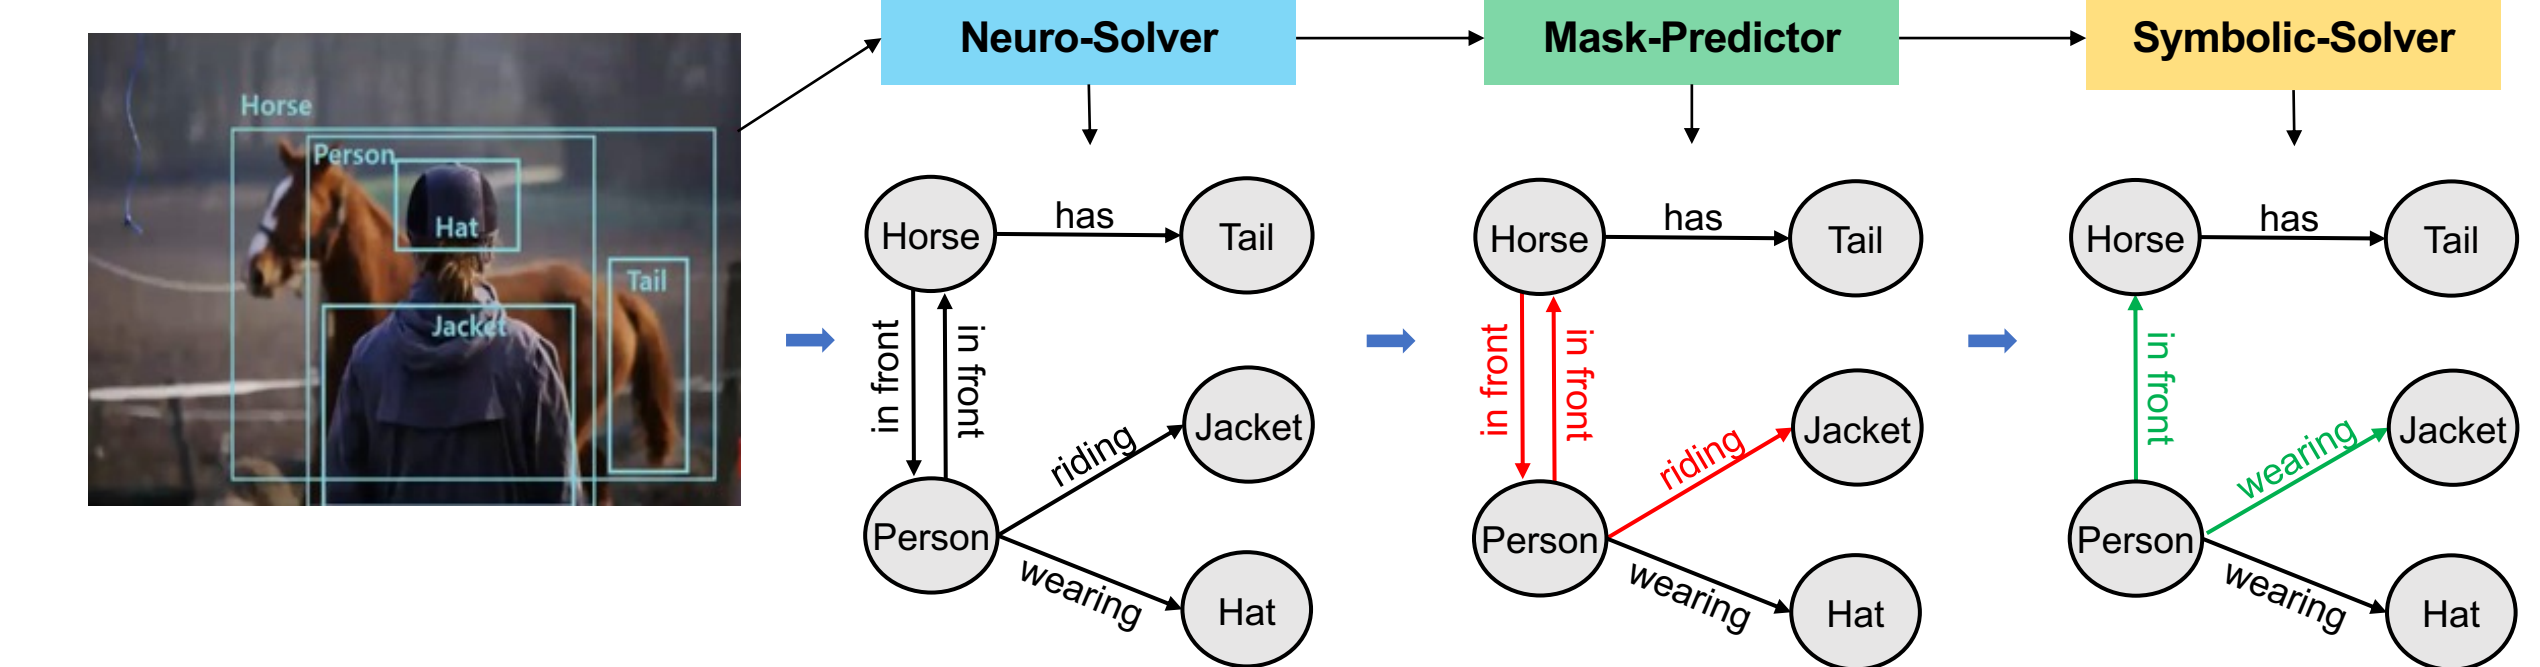
Cell(5,4)

Cell(8,8)

Neuro-Solver

It is learning the correct Sudoku rules!

Results – Scene Graph (GQA)



Task

Predicate classification

- **input**: objects ground-truth bounding boxes and labels
- **output**: scene graph

Dataset

- GQA (balanced version of VG)

Rules

- Domain/range of the predicates (e.g., $domain(wear)=\{person\}$)

Result

- Between 1% to 2% improvement



		R@20	R@50	R@100	R@200	R@300
All-shots	Baseline	29.22	42.35	48.48	50.75	51.11
	Max-improvement (PSB)	0.12	0.23	0.32	0.35	0.36
	% improvement of NASR	99.71	99.58	99.69	99.64	99.64
Zero-shots	Baseline	16.62	27.65	34.10	37.41	38.11
	Max-improvement (PSB)	0.91	1.43	1.93	2.18	2.33
	% improvement of NASR	100.00	100.00	100.00	100.00	100.00

NASR results: percentage of the max achievable improvement under the given ontology, defined by the Probabilistic Symbolic Baseline (PSB)

Take-away message

NASR: a neuro-symbolic method to manage the trade-off between the cost, expressivity, and exactness of reasoning during inference.

- Any type of rules/constraint can be used
- Results on Visual-Sudoku & Scene-Graph:
 - NASR outperforms the baseline
 - An existing method is improved when integrated in NASR
 - NASR is more efficient
 - NASR is more robust to noise compared to the symbolic baseline.



Open-Source GIT repository:

<https://github.com/corneliocristina/NASR>